



Deploying ML Models from Scratch

Ido Shlomo - Senior Data Science Manager

BlueVine

About BlueVine

- Fintech startup up based in Redwood City, CA
- Provides working capital (loans) to small & medium sized businesses
- Over \$2 BN funded to date
- Data Science challenges:
 - Consume many different semi structured / unstructured data sources
 - Deal with noisy / weak signals
 - Make decisions fast
 - Build models that are stable and accurate

About Me


- Data Science Manager @ BlueVine
- Lead BlueVine's DS team in Redwood City, CA (total of ~20 people across RWC & TLV)
- Team focus:
 - NLP & text mining
 - Anomaly detection
 - Probabilistic ML
 - Response modeling
- Personal interests: Unstructured data and DS Infrastructure.

Code & Slides

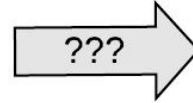
- Git Repo: <https://github.com/ido-sh>
- Slides: repo > public_presentations > [dsgo_sagemaker_2019.pdf](#)
- Code: repo > public_presentations > [sagemaker_tutorial](#)

The Issue

- In a Business context:
Be able to "deliver" something
- In Research context:
Free it from constraints of a local machine



ML Model / Code / Data

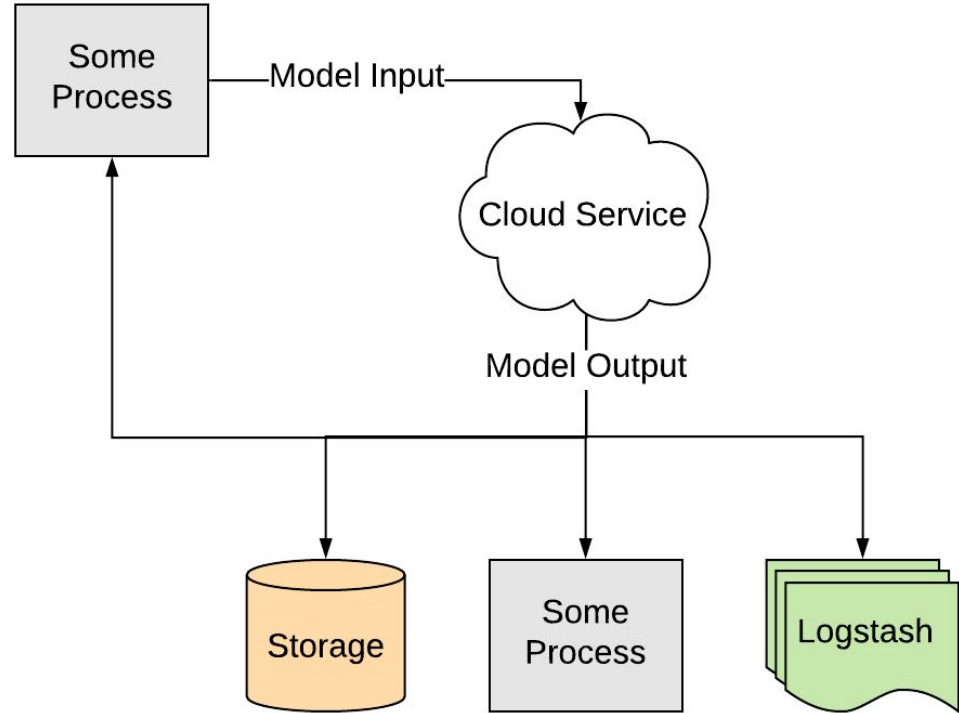


Cloud Service

The Desired End Result

Cloud Service:

- Model hosted in the cloud
- Always up (persistence)
- Communicates via API
- Can scale



The Starting Point



Desired Solution

Something that deploys our code as a cloud service

AND

In way that is:

- **Flexible:** Handle any type of Python code or structure
- **Simple:** Requires minimal effort to run
- **Independent:** Can run end-to-end by a Data Scientist with normal skill-set

Some leading deployment paradigms

Solution	Example Setups	Problem
Hand off everything to a team of Engineers		

Some leading deployment paradigms

Solution	Example Setups	Problem
Hand off everything to a team of Engineers	<ul style="list-style-type: none">● Data Scientist sends to Engineer:<ul style="list-style-type: none">○ Code & binaries○ Environment & tests○ Deployment config● Engineer:<ul style="list-style-type: none">○ Rewrites code○ Checks tests○ Deploys (somehow)	

Some leading deployment paradigms

Solution	Example Setups	Problem
Hand off everything to a team of Engineers	<ul style="list-style-type: none">• Data Scientist sends to Engineer:<ul style="list-style-type: none">○ Code & binaries○ Environment & tests○ Deployment config• Engineer:<ul style="list-style-type: none">○ Rewrites code○ Checks tests○ Deploys (somehow)	Not independent

Some leading deployment paradigms

Solution	Example Setups	Problem
Jointly manage deployment environment with Engineers		

Some leading deployment paradigms

Solution	Example Setups	Problem
Jointly manage deployment environment with Engineers	<ul style="list-style-type: none">● Data Scientist:<ul style="list-style-type: none">○ Pushes code & binaries to repo / storage○ Adheres to preset deployment config (environment, tests)● Engineer:<ul style="list-style-type: none">○ QA for code repo / storage○ Deploys (somehow)	

Some leading deployment paradigms

Solution	Example Setups	Problem
Jointly manage deployment environment with Engineers	<ul style="list-style-type: none">• Data Scientist:<ul style="list-style-type: none">○ Pushes code & binaries to repo / storage○ Adheres to preset deployment config (environment, tests)• Engineer:<ul style="list-style-type: none">○ QA for code repo / storage○ Deploys (somehow)	Semi-independent, Semi-flexible

Some leading deployment paradigms

Solution	Example Setups	Problem
Develop on a fully managed deployment-capable platform		

Some leading deployment paradigms

Solution	Example Setups	Problem
Develop on a fully managed deployment-capable platform	<ul style="list-style-type: none">• DataRobot• Alteryx• ML-Flow• Many others...	

Some leading deployment paradigms

Solution	Example Setups	Problem
Develop on a fully managed deployment-capable platform	<ul style="list-style-type: none">• DataRobot• Alteryx• Ayasdi• Many others...	Not flexible

Some leading deployment paradigms

Solution	Example Setups	Problem
Build your own docker containers and deploy them		

Some leading deployment paradigms

Solution	Example Setups	Problem
Build your own docker containers and deploy them	<ul style="list-style-type: none">• Docker to bundle code and build containers• A Docker registry to store them• A Docker orchestration tool to deploy them:<ul style="list-style-type: none">○ Kubernetes○ Docker Swarm○ Mesos○ Many others...	

Some leading deployment paradigms

Solution	Example Setups	Problem
Build your own docker containers and deploy them	<ul style="list-style-type: none">• Docker to bundle code and build containers• A Docker registry to store them• A Docker orchestration tool to deploy them:<ul style="list-style-type: none">○ Kubernetes○ Docker Swarm○ Mesos○ Many others...	Not simple

Amazon SageMaker

Motivation: Get all the pros of docker deployment without writing any docker code

- **Flexible:** Can work with any custom Python code, environment & data files
- **Simple:** Requires relatively basic Python code to run
- **Independent:** Can build docker containers and deploy them, both for cloud training jobs and deployment of cloud services

Supports four modes: Tag, Explore, **Train** and **Deploy**

Link: <https://aws.amazon.com/sagemaker/>

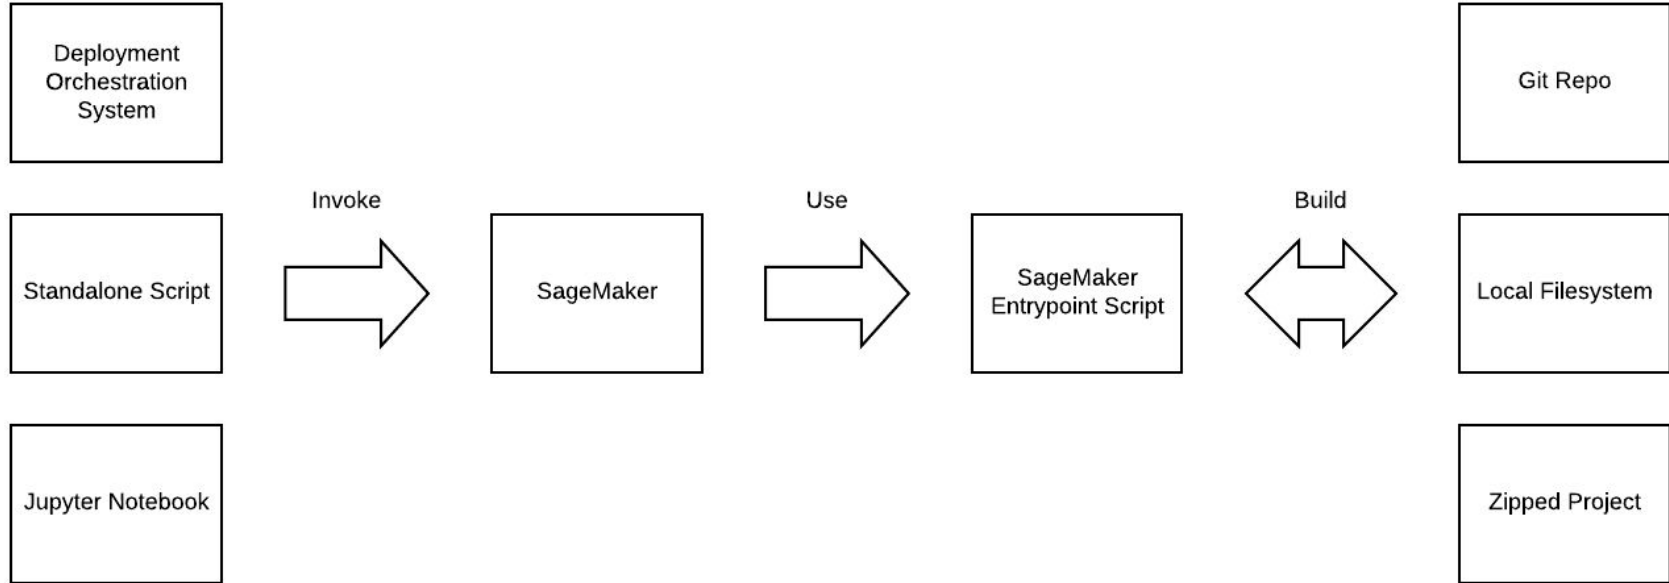
Amazon SageMaker

We're turning **ALL** of our main Data Science deliverables into cloud services

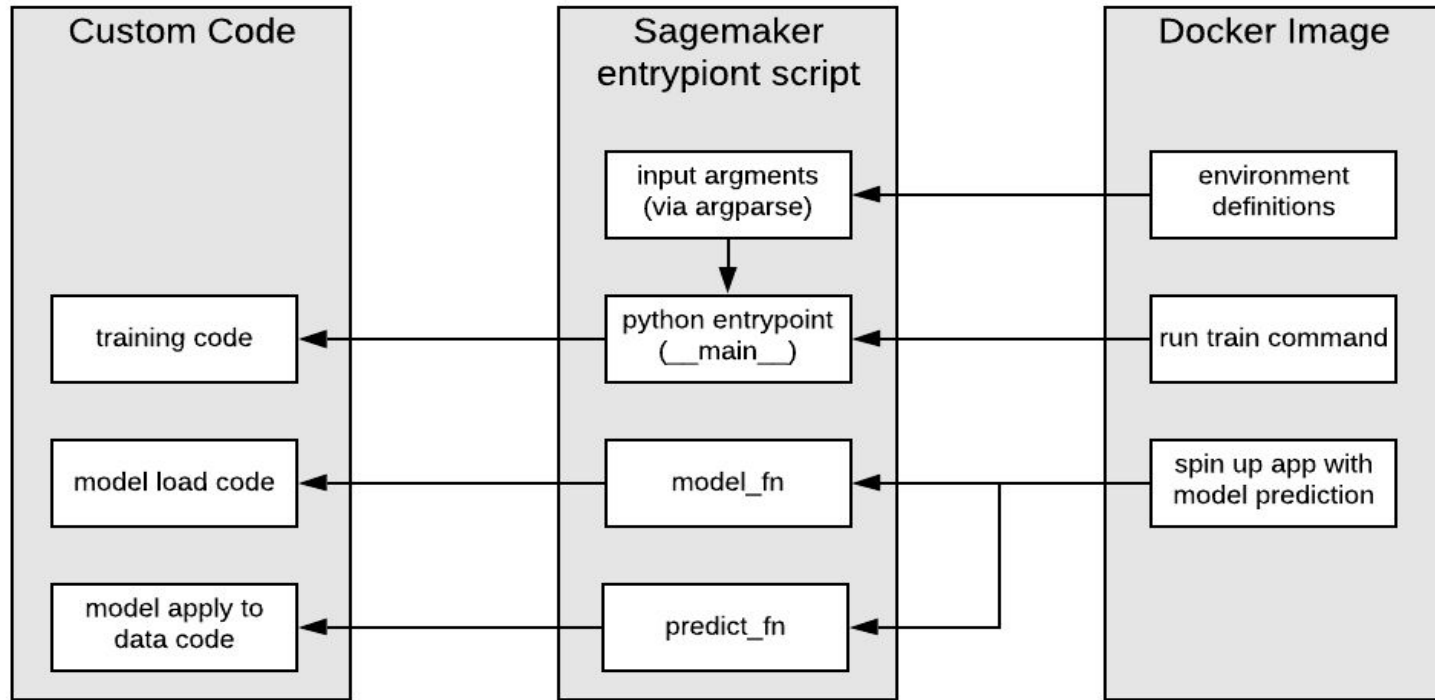
Some notable examples:

- Client Risk score
- Client Fraud score
- Industry classification
- Insights from external sources:
 - Bank data
 - Credit report data
 - Government filing data

Project Structure



Internal Connections



Example > entrypoint.py

```
import custom_code # importing whatever custom code you have

# THIS HANDLES TRAINING (DEFAULT SCRIPT INVOKE)
if __name__ == '__main__': ...

# THIS LOADS A TRAINED MODEL
def model_fn(model_dir): ...

# THIS APPLIES MODEL TO INPUT AND RETURNS PREDICTION
def predict_fn(input_data, model): ...
```

Example > entrypoint.py > __main__

```
# THIS HANDLES TRAINING (DEFAULT SCRIPT INVOKE)
if __name__ == '__main__':

    # parse environment variables
    parser = argparse.ArgumentParser()
    parser.add_argument('--output-data-dir', type=str, default=os.environ['SM_OUTPUT_DATA_DIR'])
    parser.add_argument('--model-dir', type=str, default=os.environ['SM_MODEL_DIR'])
    parser.add_argument('--train', type=str, default=os.environ['SM_CHANNEL_TRAIN'])
    args = parser.parse_args()

    # read training data from train directory
    input_files = [os.path.join(args.train, file) for file in os.listdir(args.train)]
    raw_data = [pd.read_csv(file) for file in input_files]
    train_data = pd.concat(raw_data)

    # fit model to data
    name_comparison_model = custom_code.fit_model(train_data)

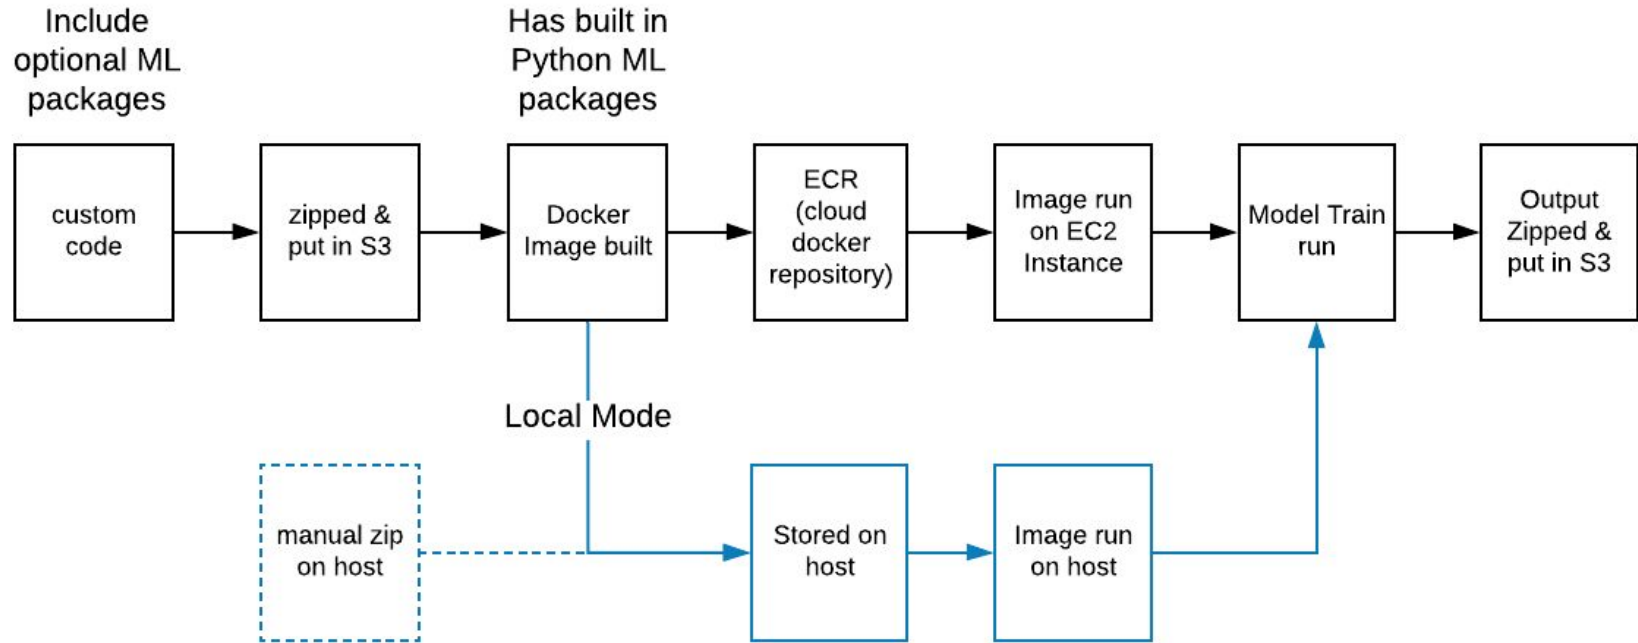
    # save model as file
    custom_code.save_model(name_comparison_model, args.model_dir)
```

Example > entrypoint.py > model_fn / input_fn / predict_fn

```
# THIS LOADS A TRAINED MODEL
def model_fn(model_dir):
    mdl = custom_code.load_model(model_dir)
    return mdl

# THIS APPLIES MODEL TO INPUT AND RETURNS PREDICTION
def predict_fn(input_data, model):
    mdl_output = custom_code.use_model(input_data, model)
    return mdl_output
```

Train Run



Train Run

```
In [1]: import sagemaker
        from boto3 import Session as BotoSession
        from time import sleep

        # create sagemaker session
        boto_session = BotoSession(profile_name='default',
                                   region_name='us-east-1')
        sagemaker_session = sagemaker.Session(boto_session=boto_session)
        sagemaker_role = 'ido-sagemaker-test'
```

Train Run

Example Model:

- Dataset of ~50K book titles from Project Gutenberg
(<https://github.com/niderhoff/nlp-datasets>)
- Train TF-IDF Vectorizer on those titles
- Build model for comparing book titles (cosine similarity)

	book_title
0	a plea for monogamy by wilfrid lay
1	jerusalem the city of herod and saladin by wal...
2	starland by robert stawell ball
3	saratoga national historical park junior range...
4	the bakhtyar nama by anonymous
5	my disillusionment in russia by emma goldman
6	the american missionary vol 37 no 2 february 1...
7	agriculture of the hidatsa indians by gilbert ...
8	the germ growers by robert potter
9	clipped wings by percy f westerman

Train Run

```
def fit_model(train_data):  
    train_text_corpus = [x for x in train_data['book_title'].tolist() if pd.notnull(x)]  
    book_title_model = TfidfVectorizer()  
    book_title_model = book_title_model.fit(train_text_corpus)  
    return book_title_model
```

```
def use_model(input_data, model):  
    input_1 = input_data['arg1']  
    input_2 = input_data['arg2']  
    logger.info('input 1: {}, input 2: {}'.format(input_1, input_2))  
    score = 1 - cosine(model.transform([input_1]).todense(),  
                        model.transform([input_2]).todense())  
    return score
```

```
def load_model(model_dir):  
    mdl = joblib.load(os.path.join(model_dir, "book_title_model.joblib"))  
    return mdl
```


__main__

predict_fn

model_fn

Train Run

```
In [2]: # upload training data to s3
train_dir = 'data/train'
project_name = 'sagemaker-dsgo-tutorial'
train_input = sagemaker_session.upload_data(
    train_dir, key_prefix="{}{}".format(project_name, train_dir))
print('location in s3: {}'.format(train_input))

location in s3: s3://sagemaker-us-east-1-/sagemaker-dsgo-tutorial/data/train
```


Train Run

```
In [3]: from sagemaker.sklearn.estimator import SKLearn

# config model training
cloud_model = SKLearn(
    entry_point='sagemaker_entry_point.py',
    source_dir='.',
    train_instance_type='ml.c4.xlarge',
    train_instance_count=1,
    role=sagemaker_role
)
```

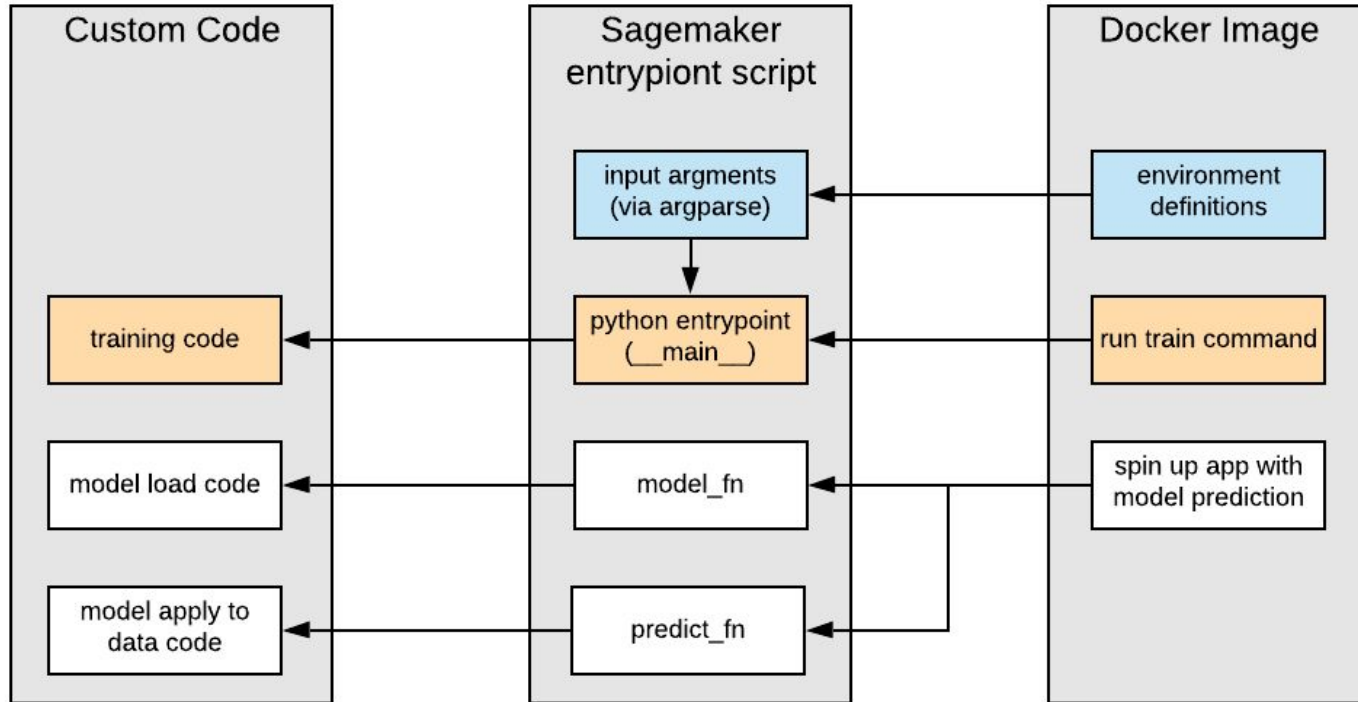
Train Run

```
In [4]: # run model training (data has to be from s3)
cloud model.fit({'train': train input})
```

Training Env:

```
{
  "input_config_dir": "/opt/ml/input/config",
  "job_name": "sagemaker-scikit-learn-XXXXXXXXXXXX",
  "module_dir": "s3://sagemaker-us-east-1-XXXXXXXXXXXX/sagemaker-scikit-learn-XXXXXXXXXXXX/source/sourcedir.tar.gz",
  "user_entry_point": "sagemaker_entry_point.py",
  "is_master": true,
  "input_dir": "/opt/ml/input",
  "log_level": 20,
  "input_data_config": {
    "train": {
      "S3DistributionType": "FullyReplicated",
      "RecordWrapperType": "None",
      "TrainingInputMode": "File"
    }
  },
  "output_dir": "/opt/ml/output",
}
```

Internal Connections



Search

▼ Ground Truth

Labeling jobs

Labeling datasets

Labeling workforces

▼ Notebook

Notebook instances

Lifecycle configurations

Git repositories

Training Panel

▼ Training

Algorithms

Training jobs

Hyperparameter tuning jobs

▼ Inference

Compilation jobs

Model packages

Models

Endpoint configurations

Endpoints

Batch transform jobs

Amazon SageMaker > Training jobs

Training jobs

Actions ▾

Create training job

🔍 Search training jobs

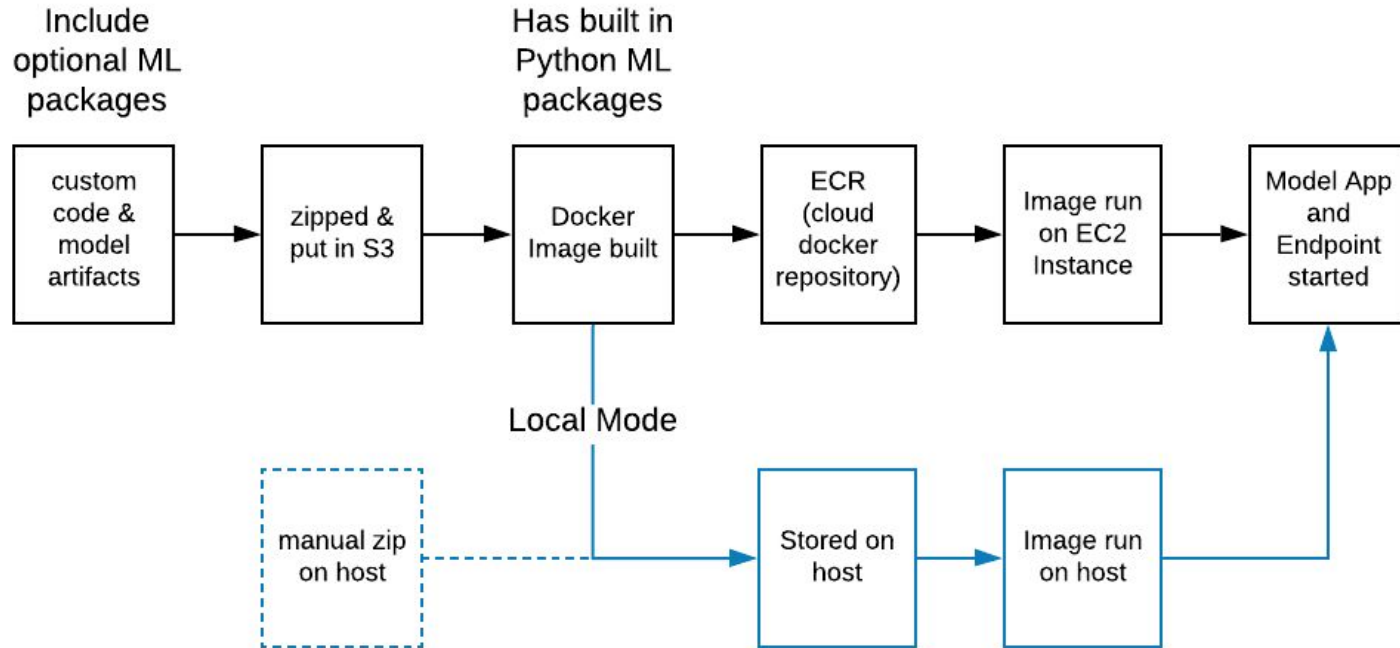


Training Jobs Panel

< 1 > ⚙️

	Name	Creation time	Duration	Status
<input type="radio"/>	sagemaker-scikit-learn-████████████████████	Sep 20, 2019 19:19 UTC	3 minutes	✅ Completed
<input type="radio"/>	sagemaker-scikit-learn-████████████████████	Sep 20, 2019 01:28 UTC	3 minutes	✅ Completed
<input type="radio"/>	sagemaker-scikit-learn-████████████████████	Sep 20, 2019 00:40 UTC	3 minutes	✅ Completed
<input type="radio"/>	sagemaker-scikit-learn-████████████████████	Sep 20, 2019 00:31 UTC	3 minutes	✅ Completed
<input type="radio"/>	sagemaker-scikit-learn-████████████████████	Sep 20, 2019 00:05 UTC	4 minutes	❌ Failed
<input type="radio"/>	sagemaker-scikit-learn-████████████████████	Sep 19, 2019 23:51 UTC	3 minutes	❌ Failed
<input type="radio"/>	sagemaker-scikit-learn-████████████████████	Sep 10, 2019 00:22 UTC	8 minutes	✅ Completed
<input type="radio"/>	sagemaker-scikit-learn-████████████████████	Sep 10, 2019 00:13 UTC	3 minutes	✅ Completed

Deploy Run



Deploy Run

```
In [5]: # from model trained on cloud via sagemaker  
cloud_predictor = cloud_model.deploy(initial_instance_count=1,  
                                     instance_type="ml.m4.xlarge")
```

-----!

Deploy Run

```
In [11]: book1 = 'tale of two cities'
book2 = 'tale by two cities'
result = cloud_predictor.predict({'arg1': book1, 'arg2': book2})

print("\nRESULT --> {} VS {}: {}".format(book1, book2, result))
```

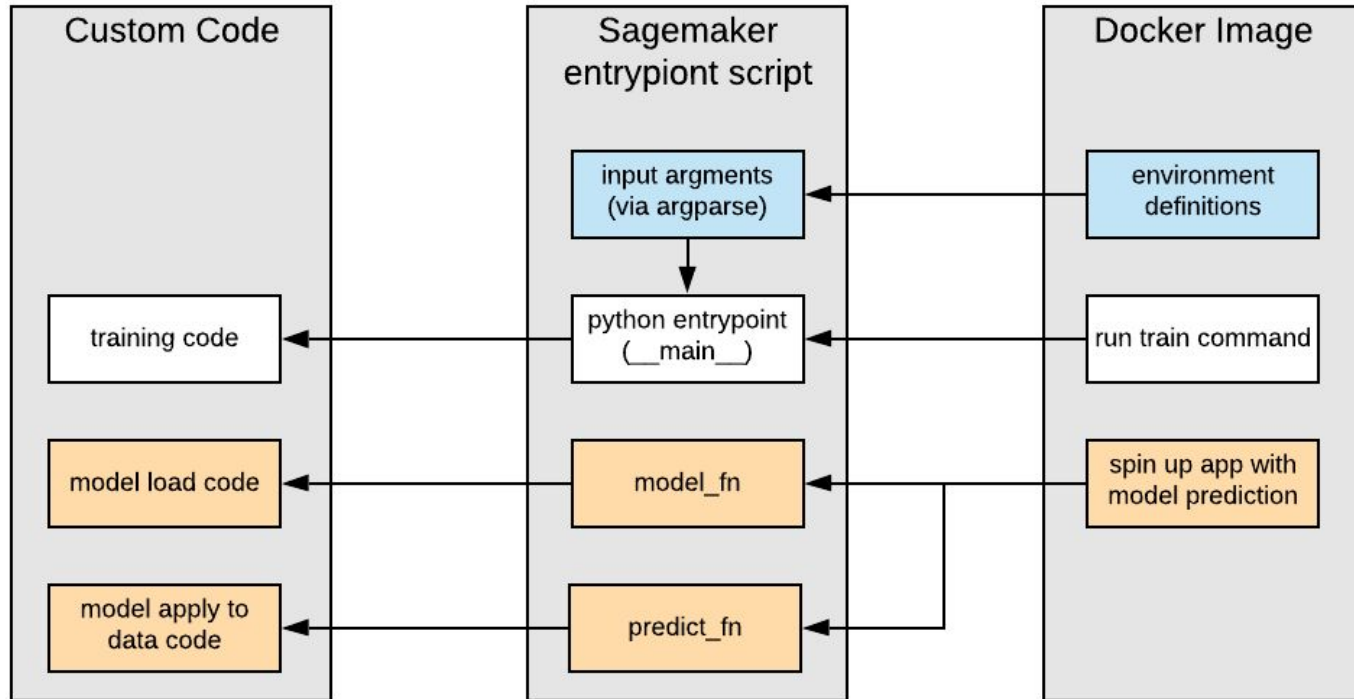
```
RESULT --> tale of two cities VS tale by two cities: 0.9797165873399724
```

```
In [12]: book1 = 'tale of two cities'
book2 = 'tale of two towns'
result = cloud_predictor.predict({'arg1': book1, 'arg2': book2})

print("\nRESULT --> {} VS {}: {}".format(book1, book2, result))
```

```
RESULT --> tale of two cities VS tale of two towns: 0.574831650820277
```

Internal Connections



▼ Ground Truth

Labeling jobs

Labeling datasets

Labeling workforces

▼ Notebook

Notebook instances

Lifecycle configurations

Git repositories

▼ Training

Algorithms

Training jobs

Hyperparameter tuning jobs

▼ Inference

Compilation jobs

Model packages

Models

Endpoint configurations

Endpoints

Batch transform jobs

Amazon SageMaker > Endpoints

Endpoints

Update endpoint

Actions ▾

Create endpoint

🔍 Search endpoints

< 1 > ⚙️



Endpoints Panel

	Name ▾	ARN	Creation time ▾	Status ▾	Last updated
<input type="radio"/>	sagemaker-scikit-learn-2019-09-26-06-14-21-532	arn:aws:sagemaker:us-east-1:██████████/sagemaker-scikit-learn-2019-09-26-06-14-21-532	Sep 26, 2019 06:22 UTC	✓ InService	Sep 26, 2019 06:29 UTC



Deployments Panel

Filter events

Time (UTC +00:00)

Message



My custom logs

2019-09-26

- ▶ 06:33:12 10.32.0.1 - - [26/Sep/2019:06:33:12 +0000] "GET /ping HTTP/1.1" 200 0 "-" "AHC/2.0"
- ▶ 06:33:17 10.32.0.1 - - [26/Sep/2019:06:33:17 +0000] "GET /ping HTTP/1.1" 200 0 "-" "AHC/2.0"
- ▼ 06:33:19 2019-09-26 06:33:19,207 INFO - custom_code - input 1: tale of two cities, input 2: tale by two cities
2019-09-26 06:33:19,207 INFO - custom_code - input 1: tale of two cities, input 2: tale by two cities
- ▼ 06:33:19 10.32.0.1 - - [26/Sep/2019:06:33:19 +0000] "POST /invocations HTTP/1.1" 200 136 "-" "AHC/2.0"
10.32.0.1 - - [26/Sep/2019:06:33:19 +0000] "POST /invocations HTTP/1.1" 200 136 "-" "AHC/2.0"
- ▶ 06:33:22 10.32.0.1 - - [26/Sep/2019:06:33:22 +0000] "GET /ping HTTP/1.1" 200 0 "-" "AHC/2.0"
- ▶ 06:33:27 10.32.0.1 - - [26/Sep/2019:06:33:27 +0000] "GET /ping HTTP/1.1" 200 0 "-" "AHC/2.0"
- ▶ 06:33:32 10.32.0.1 - - [26/Sep/2019:06:33:32 +0000] "GET /ping HTTP/1.1" 200 0 "-" "AHC/2.0"
- ▶ 06:33:37 10.32.0.1 - - [26/Sep/2019:06:33:37 +0000] "GET /ping HTTP/1.1" 200 0 "-" "AHC/2.0"
- ▶ 06:33:42 10.32.0.1 - - [26/Sep/2019:06:33:42 +0000] "GET /ping HTTP/1.1" 200 0 "-" "AHC/2.0"
- ▶ 06:33:47 10.32.0.1 - - [26/Sep/2019:06:33:47 +0000] "GET /ping HTTP/1.1" 200 0 "-" "AHC/2.0"
- ▶ 06:33:52 10.32.0.1 - - [26/Sep/2019:06:33:52 +0000] "GET /ping HTTP/1.1" 200 0 "-" "AHC/2.0"

Other Train / Deploy modes

Local Mode

- Uses docker on local machine instead of on EC2 instance in the cloud
- Train: Trained on local machine
- Deploy: Deploy an endpoint on local machine
- Does **NOT** mean offline
- Useful for debugging (avoids spin up latency & cloud computing costs)

External Model Deployment: Run deployment cycle using model trained outside SageMaker.

[Additional code examples in my repo]

Conclusion

Key Takeaways:

- To make your models “actionable”, you need to be able to deploy them
- Having a flexible, simple and independent deployment mechanism is hugely empowering
- Amazon SageMaker is one such mechanism

Some SageMaker caveats:

- Not the easiest to debug
- Local mode is not 100% offline
- Vendor lock in (Amazon)

Code & Slides

- Git Repo: <https://github.com/ido-sh>
- Slides: repo > public_presentations > [dsgo_sagemaker_2019.pdf](#)
- Code: repo > public_presentations > sagemaker_demo



DATASCIENCE **GO**
Conference 2019

We're Hiring!!!

<https://jobs.lever.co/bluevine>

- Junior Data Scientist
- Data Scientist
- Senior Data Scientist

Contact Me

- Git: <https://github.com/ido-sh>
- Email: ido.shlomo@bluevine.com

Thanks!